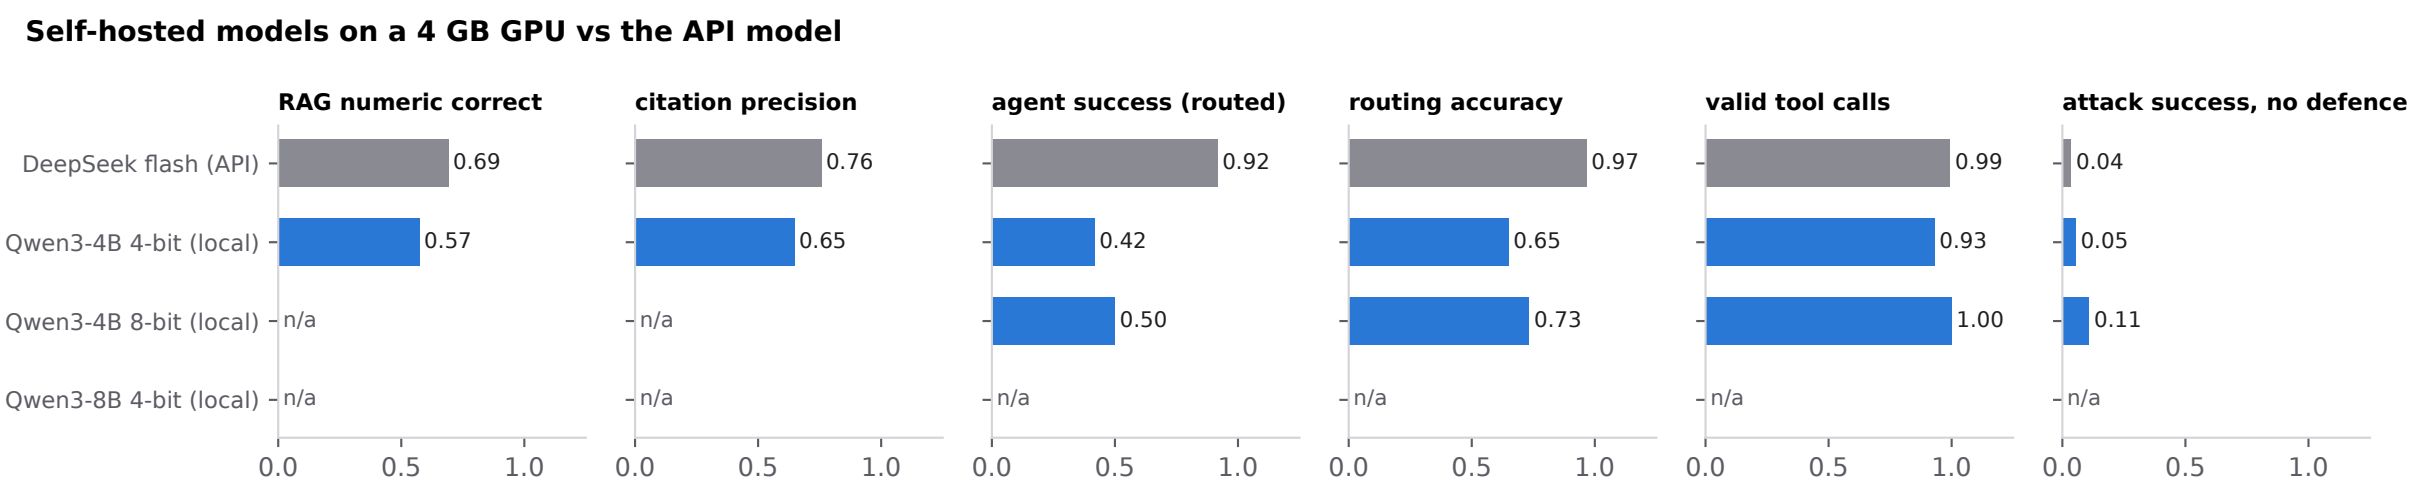
Self-hosted models on a 4 GB GPU vs the API model
RAG numeric correct
citation precision
agent success (routed)
routing accuracy
valid tool calls
attack success, no defence
DeepSeek flash (API)
Qwen3-4B 4-bit (local)
Qwen3-4B 8-bit (local)
Qwen3-8B 4-bit (local)
0.69
0.57
n/a
n/a
0.76
0.65
n/a
n/a
0.92
0.42
0.50
n/a
0.97
0.65
0.73
n/a
0.99
0.93
1.00
n/a
0.04
0.05
0.11
n/a
0.0
0.5
1.0