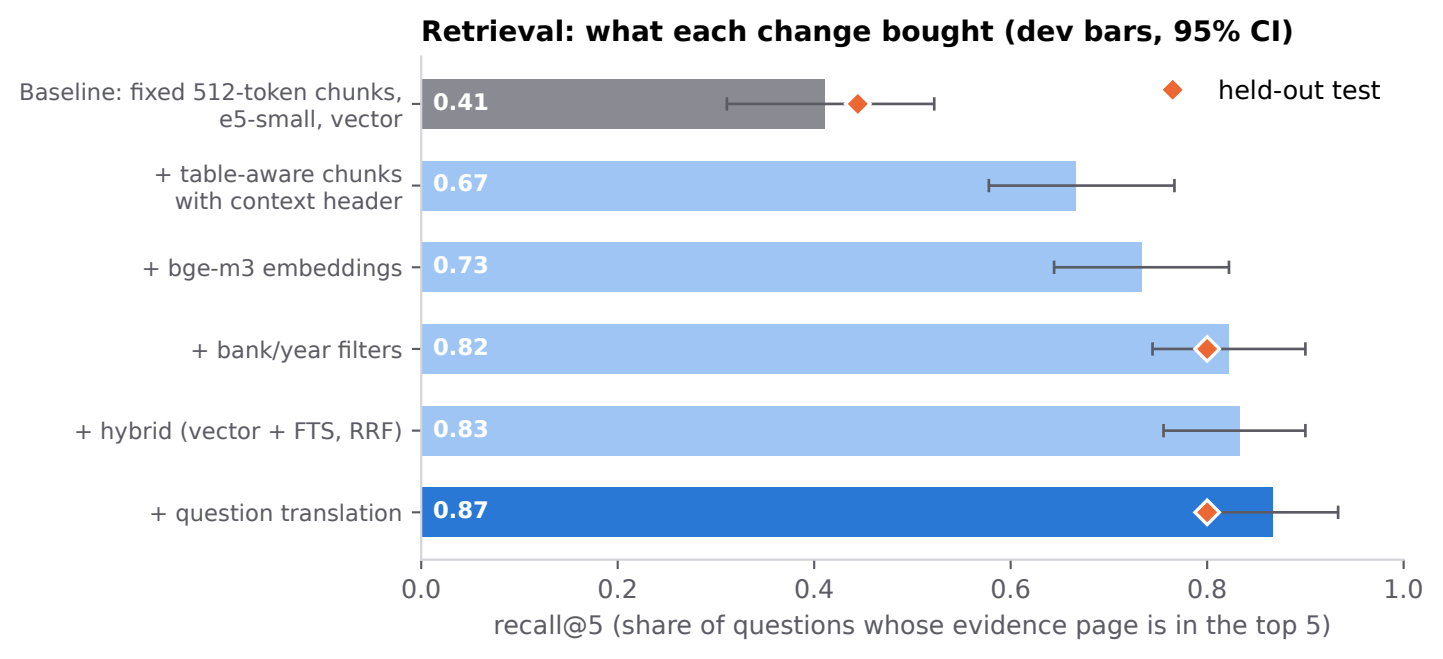
Retrieval: what each change bought (dev bars, 95% CI)
held-out test
Baseline: fixed 512-token chunks, e5-small, vector
0.41
+ table-aware chunks with context header
0.67
+ bge-m3 embeddings
0.73
+ bank/year filters
0.82
+ hybrid (vector + FTS, RRF)
0.83
+ question translation
0.87
0.0
0.2
0.4
0.6
0.8
1.0
recall@5 (share of questions whose evidence page is in the top 5)